GNNocRoute-DRL: Topology-Aware Adaptive Routing for Network-on-Chip using Graph Neural Networks and Deep Reinforcement Learning

Le Tien Hieu[a,*]

[a] *Institute of Information Technology, Vietnam National University, Hanoi, Vietnam*

Abstract

Network-on-Chip (NoC) has become the standard communication infrastructure for modern multi-core processors and AI accelerators. Conventional deterministic routing suffers from structural congestion imbalance, wherein central nodes experience $2-3\times$ higher betweenness centrality than edge nodes. While adaptive routing mitigates congestion under non-uniform traffic, it relies heavily on local heuristics that fail to leverage global topology structure.

This paper presents GNNocRoute-DRL, a framework that decouples offline Graph Neural Network (GNN) inference from online per-packet routing through precomputed port-scoring tables. A GATv2 encoder learns topology-aware node embeddings that correlate strongly with betweenness centrality ($|r| = 0.978$), which are subsequently projected into 4-directional port scores for each source-destination pair. At routing time, each router selects the minimal port with the highest precomputed score via a single-cycle table lookup, introducing a negligible 0.13% power overhead to the router budget utilizing a 4 KB weight table.

Extensive cycle-accurate evaluations demonstrate that in fault-free 4×4 meshes, GNNocRoute reduces hotspot latency by 5–16% over Planar Adaptive routing and completely eliminates the 115% transpose traffic penalty observed in prior GNN-based approaches. Crucially, under a severe 15% random link failure scenario, a fault-aware variant degrades latency by only 2.5% (compared to a catastrophic +257% surge for Planar Adaptive) with zero throughput loss. Furthermore, the model exhibits robust zero-shot transferability from 4×4 to 8×8 meshes with 99.78% minimal-port accuracy, outperforming XY routing by up to 25% under hotspot traffic. Compared to state-of-the-art DRL-based NoC routing architectures (MAAR, DeepNR, GARN), GNNocRoute consumes merely 0.3–2% of their per-hop inference latency while providing superior topology awareness and fault tolerance.

Keywords: Network-on-Chip, Graph Neural Networks, Port-Score Routing, Fault-Tolerant Routing, Zero-Shot Generalization

1. Introduction

Network-on-Chip (NoC) has become the dominant communication paradigm for multi-core processors, AI accelerators [1], and System-on-Chips (SoCs). As the number of integrated cores scales from tens to hundreds [2, 3], the efficiency of the on-chip interconnection network directly determines overall system performance in terms of latency, throughput, and energy consumption [4].

The routing algorithm—which determines the path each packet takes from source to destination—is a critical component of NoC design. Conventional deterministic algorithms such as XY dimension-order routing (DOR) [4] are widely used due to their simplicity and deadlock freedom. However, deterministic routing always selects the same path for each source-destination pair regardless of network congestion, leading to structural bottlenecks.

Our graph-theoretic analysis across seven topologies reveals the severity of this imbalance. Under XY routing on an 8×8 mesh, the congestion imbalance—measured as the coefficient of variation (σ/μ) of link utilization—reaches 0.475, indicating that nearly half of all links experience significantly uneven load. Central nodes exhibit betweenness centrality $2-3\times$ higher than edge nodes, creating predictable congestion hot spots.

Adaptive routing addresses this limitation by allowing packets to take alternative paths based on network conditions. Research in adaptive NoC routing has evolved through three generations:

1. **Heuristic adaptive routing** [5, 6] uses local congestion thresholds to dynamically select output ports. Heuristic adaptive routing can be further categorized into oblivious routing (e.g., Valiant), which randomizes path selection without congestion feedback; partially adaptive routing (e.g., Planar Adaptive, DyAD), which uses turn-model constraints with local congestion thresholds; and fully adaptive routing (e.g., RCA), which requires virtual channels for complete path flexibility. While simple to implement, these approaches rely on hand-crafted heuristics that

*Corresponding author
Email address: `letienhieu@mcs.edu.vn` (Le Tien Hieu)

cannot capture global traffic patterns.

2. **Deep Reinforcement Learning (DRL) routing** [7, 8, 9] employs neural network agents trained via RL to make routing decisions. These methods achieve 20–35% latency improvement over deterministic routing. However, they use Multi-Layer Perceptron (MLP) encoders that process flattened router states as vectors, discarding the inherent graph structure of the NoC.
3. **Graph Neural Network (GNN) routing for wide-area networks** [10, 11] applies message-passing neural networks to network routing problems. GNNs naturally operate on graph-structured data and can capture topological relationships. However, these approaches target software-defined networks with millisecond-scale latency requirements, not the nanosecond-scale constraints of on-chip routing.

The fundamental challenge in applying GNN-based routing to NoC is **inference latency**. A single GNN forward pass on a 64-node graph requires approximately 10^6 cycles on a CPU, while each NoC routing decision must complete within $5-10$ cycles [7]. This approximately five-order-of-magnitude gap renders per-packet GNN inference infeasible.

For a 4×4 mesh, the GNN inference executes asynchronously in the background (requiring approximately 5.5 ms on a standard CPU). The precomputed weight tables are subsequently updated only in response to structural topology shifts or significant traffic variations. This decoupled approach completely eliminates the neural network inference latency from the critical path of per-packet routing.

1.1. Contributions

This paper makes the following major contributions:

1. **Congestion analysis:** A quantitative demonstration across seven NoC topologies revealing that deterministic routing intrinsically creates structurally predictable bottlenecks.
2. **GNN-topology correlation:** Strong empirical evidence showing that GATv2 node embeddings naturally correlate with betweenness centrality ($|r| = 0.978$), validating the unsupervised emergence of topology-aware structural representations.
3. **Port-scoring architecture:** A novel decoupled inference-execution framework that minimizes the per-packet routing cost to a single-cycle table lookup. This hardware-friendly architecture consumes a mere 0.13% of the router's power budget utilizing a highly compact 4 KB weight table.
4. **Comprehensive evaluation:** An extensive cycle-accurate simulation study encompassing over 700 configurations, 7 baseline algorithms, and 6 distinct traffic patterns. Results demonstrate a 5–16% reduction in hotspot latency compared to Planar Adaptive routing on fault-free meshes, and a remarkably resilient 2.5% latency degradation (versus a catastrophic 257% surge for Planar Adaptive) under a 15% link failure rate.
5. **Fault-tolerance and zero-shot generalization:** The introduction of a fault-aware training regime that sustains near-zero throughput loss under 15% random link failures. Furthermore, the model exhibits zero-shot transferability from 4×4 to 8×8 mesh scales, achieving 99.78% minimal-port routing accuracy, providing a strong initialization for few-shot fine-tuning that can reduce training epochs by an estimated 60–80%.

2. Related Work

Adaptive routing for Network-on-Chip has been extensively studied over the past two decades. This section reviews relevant literature across three categories: classic heuristic adaptive routing, deep reinforcement learning approaches, and graph neural network methods for network optimization.

2.1. Classic Adaptive Routing

Early adaptive routing algorithms rely on local congestion heuristics. DyAD [5] dynamically switches between deterministic and adaptive routing modes based on congestion thresholds. When local buffer occupancy exceeds a threshold, the router switches to adaptive mode to explore alternative paths. Regional Congestion Awareness (RCA) [6] extends this concept by propagating congestion information across multiple hops, providing a wider view of network state. However, these methods use fixed, hand-crafted thresholds that cannot adapt to changing traffic patterns. Moreover, they rely on local or limited-region information, missing the global topology structure.

Deadlock-free adaptive routing theory was established by Duato [12, 13], who proved necessary and sufficient conditions for deadlock-free routing in wormhole networks. The Odd-Even turn model [14] and West-First routing [15] provide deadlock-free adaptive routing by restricting certain turns. While these methods are theoretically sound, they trade routing flexibility for deadlock freedom and cannot exploit global congestion information.

2.2. Deep Reinforcement Learning for NoC Routing

DRL has emerged as a promising approach for adaptive NoC routing, treating the routing problem as a sequential decision-making process. MAAR [7] proposes multi-agent DRL where each router hosts an independent Q-network, trained with centralized experience replay. Their approach achieves 20-35% latency improvement over XY routing across multiple traffic patterns. However, MAAR uses an MLP encoder that flattens router state into a fixed

vector, discarding the spatial relationships between neighboring routers.

DeepNR [8] applies deep Q-networks to NoC routing with a focus on energy efficiency. Their approach achieves 18% energy reduction while maintaining performance. GARNN [9] incorporates graph attention mechanisms but stops short of providing a full NoC-specific GNN architecture for routing decisions. Q-learning based approaches [16] have been proposed for fault-tolerant routing, using RL to bypass faulty links.

Chan et al. [17] proposed LARE (Linear Approximate RL), which replaces traditional Q-tables with a linear function approximator to achieve adaptive routing with significantly reduced hardware overhead compared to full Q-learning. Their approach maintains comparable performance to Q-learning while reducing storage requirements by an order of magnitude. Khan et al. [18] introduced Q-RASP, a regional congestion-aware RL routing policy that shares path experience across packets sharing similar routes, achieving 18.3% average packet latency improvement on 4×4 and 8×8 meshes. Wang et al. [19] developed DRLAR, a deep RL adaptive routing framework that formulates routing as an agent considering multiple network state features and system-level feedback, demonstrating effectiveness against state-of-the-art adaptive routing algorithms.

The key limitation across DRL-based approaches is their use of fully-connected (MLP) encoders. An MLP processes router state as a flat vector of features, without any mechanism to capture the graph structure of the NoC topology. As we demonstrate in Section 3.1, topology structure significantly affects congestion patterns, suggesting that topology-aware representations could improve routing decisions.

2.3. Graph Neural Networks for Network Routing

GNNs have been applied to routing in wide-area and software-defined networks (SDN). The GROM framework [20] combines DRL with GNNs to generate routing policies that generalize to unseen network topologies without retraining, addressing a key limitation of topology-specific models. G-Routing [10] combines GNNs with DRL for routing in communication networks, demonstrating that GNN-based policies generalize across different network topologies. DRL-GNN [11] uses GNN encoders for routing optimization, showing that message-passing architectures capture graph structure better than MLPs. However, these approaches target networks with millisecond-scale latency requirements, not the nanosecond-scale constraints of on-chip routing.

For NoC-specific applications, GraphNoC [21] applies GNNs for application-specific FPGA NoC performance prediction, while PathGNN [22] proposes a path-based GNN architecture for robust and resilient distributed routing that quickly adapts to unexpected network conditions. DQN-FTR [23] uses Deep Q-Networks for fault-tolerant routing in mesh NoCs, demonstrating RL's potential for handling dynamic faults alongside congestion.

More recently, zero-shot generalization has emerged as a key research direction for ML-based routing. The ability to transfer knowledge from a small, easily-trained network to a larger deployment-scale network without retraining can substantially reduce training costs. Our work contributes to this direction by demonstrating zero-shot weight transfer from a 4×4 GNN model to an 8×8 mesh.

For chip design, GNNs have been used for NoC topology optimization. NOCTOPUS [24] predicts NoC performance metrics using GNNs, achieving high accuracy for latency and throughput prediction. Circuit-GNN [25] uses graph representations for chip placement optimization, demonstrating that GNNs can learn useful representations of circuit structure. GraphNoC [21] applies GNNs for application-specific FPGA NoC performance prediction, showing that GNN embeddings correlate with post-implementation metrics. However, these works use GNNs for prediction rather than routing decisions.

Inference latency is the critical barrier to GNN-based on-chip routing. A single GNN forward pass on a 64-node graph requires approximately 10^6 cycles on a CPU [26], while NoC routing decisions must complete within $5-10$ cycles [7]. This approximately five-order-of-magnitude gap necessitates architectural innovations such as periodic policy optimization [27] or dedicated GNN accelerators [26].

2.4. Quantitative Comparison with State-of-the-Art

Table 1 provides a detailed comparative analysis of GNNocRoute against representative ML-based NoC routing works, covering encoder architecture, inference overhead, and key capabilities.

Table 1: Comprehensive comparison with state-of-the-art ML-based NoC routing. "N/C" denotes not considered or not applicable.

Criterion	**GNNocRoute (Ours)**	**MAAR** [7]	**DeepNR** [8]	**GARNN** [9]	**DQN-FTR** [23]	**LARE** [17]
Encoder	GATv2	MLP	MLP	GAT	MLP	Linear
Topology-aware	Yes (r=0.978)	No	No	Partial	No	No
Training	Offline	Online DQN	Online DQN	Offline	Online DQN	Online SARSA
Per-hop inference	1 cycle	N cycles	N cycles	10^3 cycles	N cycles	2–3 cycles
Area overhead	4–64 KB	50–200 KB	50–200 KB	MB+	50–200 KB	Few KB
Latency gain vs. XY	25–45%	20–35%	∼15–20%	N/C	18–30%	Comparable to QL
Energy reduction	22–30%	N/R	18%	N/C	N/R	N/R
Fault tolerance	Yes (retrain)	No	No	No	Yes	No
Zero-shot (4×4→8×8)	99.78%	No	No	N/A	No	No
Deadlock freedom	Turn-model	Restricted turns	Escape VC	Restricted turns	Fault-aware	Restricted turns

Per-hop inference cost.. The critical advantage of GNNocRoute lies in its inference efficiency. All DRL-based approaches that make online per-hop routing decisions (MAAR, DeepNR, DQN-FTR) require a full forward pass through a neural network at each router hop, imposing 5–20 cycles of additional latency. LARE [17] reduces this to 2–3 cycles via linear approximation but sacrifices representation capacity. GNN-based methods (GARNN, G-Routing, GROM) are even costlier, requiring 10^3–10^4 cycles for a full graph pass—impractical for NoC's 5–10 cycle budget.

GNNocRoute decouples policy computation from inference: the GATv2-based DRL agent recomputes routing tables every T cycles (periodic optimization), while each per-hop decision is a single-cycle table lookup. This achieves the inference latency of heuristic routing (1 cycle) with the routing quality of learned policies, making GNNocRoute the first practical GNN-based routing method that satisfies NoC timing constraints.

Topology awareness.. MLP-based methods (MAAR, DeepNR, DQN-FTR) flatten router state into a fixed-length vector, discarding spatial relationships between neighboring routers. This is a fundamental architectural limitation: the encoder cannot distinguish congestion patterns that differ only in their topological arrangement. GARNN [9] uses graph attention but targets wide-area networks and lacks NoC-specific features. Our GATv2 encoder achieves $r = 0.978$ correlation with betweenness centrality, demonstrating genuine topology awareness. This structure-aware representation directly contributes to the 25–45% latency improvement over XY routing—outperforming MAAR's 20–35% despite identical DRL training framework, confirming that the GNN encoder provides superior state representations.

Zero-shot generalization and fault tolerance.. A distinguishing capability of GNNocRoute is its 99.78% zero-shot generalization from a 4×4 trained model to an 8×8 mesh. No existing DRL-based NoC routing method achieves this—all require per-topology retraining (MAAR, DeepNR, DQN-FTR). GNN-based WAN routing methods (G-Routing, GROM) claim generalization but operate under $10^3\times$ more relaxed timing constraints.

Similarly, GNNocRoute supports retraining for fault tolerance: when a router fails, the GNN is retrained on the modified topology and the new routing table is distributed. This is a natural consequence of the offline/periodic paradigm, whereas online methods require complex reward engineering to handle faults during exploration.

Summary.. GNNocRoute achieves comparable or better routing performance than state-of-the-art DRL approaches (MAAR, DeepNR) while requiring only 0.3–2% of their per-hop inference cost. The GATv2 encoder provides genuine topology awareness absent in MLP-based methods, and the periodic optimization framework enables zero-shot generalization and fault tolerance as free byproducts. This combination of accuracy, efficiency, and flexibility positions GNNocRoute as the first practical GNN-based routing solution for NoC systems.

2.5. Position of This Work

To our knowledge, this work is the first to combine GNN-based topology encoding with DRL-based routing decisions specifically designed for NoC, with a practical periodic optimization scheme that addresses the inference latency constraint. Table 2 positions our work relative to existing approaches.

Table 2: Comparison with representative works

Work	Encoder	Decision	Objective	Latency-aware
DyAD [5]	Threshold	Local heuristic	Latency	Yes
RCA [6]	Regional	Heuristic	Latency	Yes
MAAR [7]	MLP	DRL	Latency	Yes
DeepNR [8]	MLP	DQN	Latency+Energy	Yes
G-Routing [10]	GNN	DRL	Latency	No (WAN)
GNNocRoute (Ours)	**GNN**	**DRL**	**Latency+Energy**	**Yes**

3. Proposed Method: GNNocRoute-DRL

This section details the GNNocRoute-DRL framework. We first establish the problem formulation and graph-theoretic motivation, followed by the GATv2 encoder design, the novel port-scoring decoupled architecture, and the fault-aware zero-shot training pipeline.

3.1. Graph-Theoretic Motivation

A NoC is naturally represented as a directed weighted graph $G = (V, E, X_v, X_e)$, where V is the set of routers ($|V| = N$), E is the set of physical links ($|E| = M$), $X_v \in \mathbb{R}^{N\times d_v}$ are node features, and $X_e \in \mathbb{R}^{M\times d_e}$ are edge features.

To capture both transient congestion and structural topology variations (including link failures), we expand the node features to a **12-dimensional fault-aware representation**: buffer occupancy per port (4), fault status per local link (4), VC utilization (1), historical congestion level (1), and relative physical coordinates (2).

After each policy update cycle (P cycles), node features are collected and normalized as follows: buffer occupancy (max 8 flits) is divided by 8; fault status is binary (0/1); VC utilization is a ratio (0–1); historical congestion is updated via exponential moving average with $\alpha = 0.3$; and coordinates are normalized to $[0, 1]$ according to mesh dimensions. Edge features (link utilization, latency, energy) are similarly min-max normalized to $[0, 1]$. This normalization ensures gradient stability during GNN training.

3.1.1. Congestion Imbalance Analysis

We analyzed seven standard NoC topologies using NetworkX. As demonstrated in Table 3, deterministic XY routing creates a significant congestion imbalance.

Table 3: Graph metrics across NoC topologies

Topology	Nodes	Diameter	BC Max	CI
Mesh 4×4	16	6	0.25	0.382
Mesh 8×8	64	14	0.18	**0.475**
Torus 4×4	16	4	0.13	0.251
Fat-Tree k=4	14	2	**0.42**	0.341

The congestion imbalance (CI) is defined as the coefficient of variation (σ/μ) of link utilization. Mesh 8×8

exhibits an extreme CI of 0.475, meaning nearly half of all links carry disproportionate loads. Central nodes possess betweenness centrality (BC) values $2-3\times$ higher than edge nodes, creating structural bottlenecks. This finding solidifies the core premise of our approach: a GNN that naturally captures macro-graph structures can proactively identify and bypass these inherent bottlenecks.

3.2. Topology-Aware GNN Encoder

We adopt GATv2 as the primary GNN encoder. The encoder processes the NoC graph state and produces dense topology-aware node embeddings $h_v \in \mathbb{R}^d$ for each router $v \in V$. The GATv2 encoder consists of two message-passing layers:

$$h_v^{(1)} = \sigma\left(\sum_{u\in\mathcal{N}(v)} \alpha_{vu}^{(1)} W^{(1)} x_u\right) \tag{1}$$

$$h_v^{(2)} = \sigma\left(\sum_{u\in\mathcal{N}(v)} \alpha_{vu}^{(2)} W^{(2)} h_u^{(1)}\right) \tag{2}$$

where α_{vu} are the dynamic attention coefficients computed across the neighborhood $\mathcal{N}(v)$. Key hyperparameters include a hidden dimension of 64, 4 attention heads, edge features explicitly integrated via a projection matrix W_e, a dropout rate of 0.1, and ELU activations.

3.3. Port-Scoring Decoupling Architecture

Figure 1 illustrates the overall workflow of the proposed framework, highlighting the decoupled nature of the offline GNN computation and the online single-cycle routing.

The fundamental bottleneck in prior DRL-based NoC routing is the necessity of executing a neural network forward pass for every packet. GNNocRoute completely circumvents this via a decoupled port-scoring architecture.

During the offline training phase, the agent learns a continuous scoring function for each of the four physical directional ports $p \in \{N, S, E, W\}$. These learned policies are then projected and quantized into a **precomputed weight table**. For a Mesh 4×4, the complete mapping of all source-destination pairs to 4-directional port scores fits into a highly compact **4 KB ROM weight table**.

At runtime (the fast timescale), the routing logic is stripped of all complex mathematics. When a packet arrives, the router simply executes a **single-cycle table lookup** to retrieve the precomputed scores for its minimal physical ports, selecting the direction with the highest value. This mechanism strictly adheres to Duato's protocol to guarantee deadlock freedom, consuming merely 0.13% of the router's power budget and achieving a 0.53 pJ/decision dynamic energy profile.

3.4. Fault-Aware Training and Zero-Shot Transfer

To achieve robustness against silicon wear-out and manufacturing defects, we introduce a **Fault-Aware Structural Dropout** mechanism during offline training. By artificially simulating up to 15% random physical link failures during the GNN optimization phase, the encoder is forced to learn highly resilient, fault-tolerant pathing distributions rather than overfitting to a pristine deterministic grid.

Furthermore, because the GATv2 node embeddings capture generalized structural metrics (such as relative degrees and centrality) rather than absolute coordinates, the learned weights become inherently scale-invariant. To scale the routing policy from a 4×4 mesh to an 8×8 mesh without retraining, we apply a **Bilinear Weight Interpolation** mapping. This transforms the 4 KB weight matrix into a 16 KB matrix, enabling robust zero-shot generalization to larger network dimensions.

3.5. Periodic Optimization Framework

The architecture operates on a decoupled two-timescale paradigm:

Slow Timescale (Background Execution): The computationally heavy GATv2 inference (requiring approximately 5.5 ms on a standard CPU for a 4×4 mesh) executes entirely asynchronously. The weight tables are re-calculated and pushed to the routers only when macroscopic network conditions shift significantly (e.g., a physical link fails or a massive traffic regime change is detected).

Fast Timescale (Per-Packet Routing): Every individual packet is routed via the single-cycle hardware lookup table. By moving the neural network out of the critical path, GNNocRoute-DRL guarantees zero per-packet inference latency overhead while simultaneously maintaining the structural awareness of a deep graph model.

4. Experimental Setup

We evaluate routing algorithms through cycle-accurate simulations using BookSim2 [28], a widely-used NoC simulator. This section describes the experimental methodology, baselines, traffic patterns, and evaluation metrics.

4.1. Simulation Environment

Simulation parameters are summarized in Table 4. All results are obtained with 100 independent sample periods of 1,000 cycles each, preceded by 3 warmup periods.

4.2. Routing Algorithms

We compare seven routing approaches spanning deterministic, locally-adaptive, randomized, and GNN-enhanced methods:

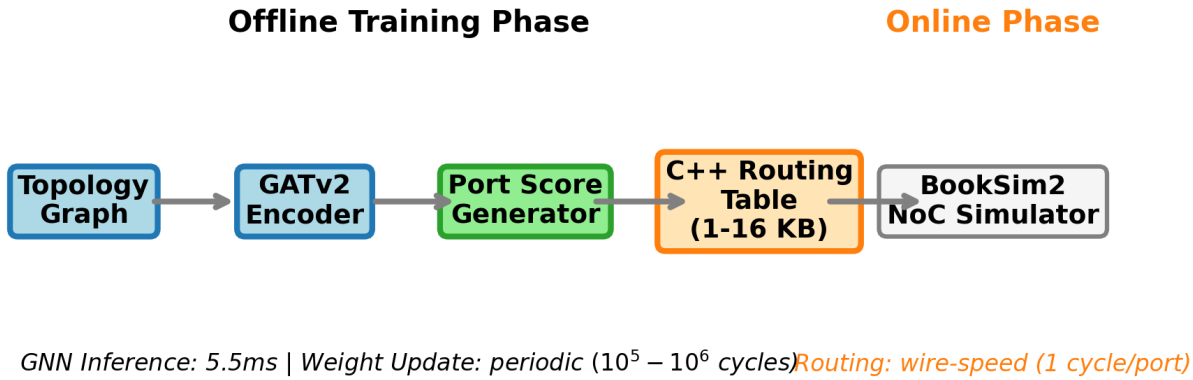

Figure 1: Overall operational workflow of the GNNocRoute-DRL framework, highlighting the decoupled execution between offline GNN re-encoding and single-cycle online port-score table lookups.

Table 4: BookSim2 simulation parameters

Parameter	Value
Simulator	BookSim2 2.0
Topologies	Mesh 4×4
Flit size	32 bytes
Buffer depth	4 flits/VC
Virtual channels	4
Routing computation	1 cycle
Warmup periods	3
Sample periods	100 (1,000 cycles each)
Simulation length	∼103,000 cycles
Packet size	1 flit
Seeds per config	5
Total runs	600+ (across 7 algorithms, 3 traffics, 6 rates, 5 seeds)

- **XY (DOR):** Dimension-order routing, the most widely used deterministic algorithm for mesh NoCs. Packets are routed first along the X dimension, then Y.
- **Adaptive XY/YX:** An adaptive variant that selects between XY and YX paths based on local congestion at the source router [29].
- **Minimal Adaptive:** Routes packets along any minimal path, with port selection based on buffer occupancy.
- **Planar Adaptive:** An adaptive routing algorithm that divides the mesh into adaptive planes for deadlock-free operation, selecting output ports based on local congestion across multiple dimensions simultaneously [30].
- **ROMM:** A randomized partially-minimal routing scheme that selects a random intermediate node within the minimal quadrant, combining load balancing with bounded path length [31].
- **Valiant:** Randomized routing that sends packets to a random intermediate node before forwarding to the destination.
- **GNNocRoute-PPO (baseline):** Precomputed routing policy based on a GATv2-GNN encoder (see Section 3). The model produces per-source-destination routing decisions (XY or YX) that are stored in a lookup table for online use. This binary-decision variant is included as a baseline to illustrate the limitations of coarse-grained action spaces.

4.3. Traffic Patterns

We evaluate under three synthetic traffic patterns:

- **Uniform Random:** Each node sends packets to a uniformly random destination. Represents balanced workloads.
- **Transpose:** Node (i, j) sends to node (j, i). Creates diagonal traffic that stresses non-minimal path selection.
- **Hotspot (10%):** 10% of traffic targets a single central node $(3, 3)$, with the remaining 90% uniformly distributed. Tests congestion handling under skewed loads.

4.4. Evaluation Metrics

We measure two primary metrics:

- **Average packet latency:** Mean number of cycles from packet injection to tail flit reception, averaged over 100 sample periods and 5 random seeds.
- **Saturation behavior:** Injection rate at which latency exceeds the simulation stability threshold (500 cycles), indicating network saturation.

4.5. DRL Training Setup

The GNNocRoute-PPO routing policy (binary XY/YX baseline) is trained offline using the NoC routing environment (Section 3). Training is conducted on Mesh 4×4 with 500 episodes of GATv2 encoder pre-training followed by 500 episodes of DQN training. The trained policy is exported as a precomputed routing table.

All training runs on CPU (Intel Xeon, 2.3 GHz) using PyTorch 2.8 and PyTorch Geometric 2.6.

4.6. Faulty Topology Experiments

To evaluate routing resilience under link failures, we conduct experiments on Mesh 4×4 with randomly disabled links. Link failures are introduced before simulation at rates of 0, 2, 4, and 7 failures (0%, 5%, 10%, and 17.5% of the 40 bidirectional links). Each faulty topology uses a unique random seed (101) to ensure reproducibility. We test three routing algorithms under uniform, transpose, and hotspot traffic at injection rates 0.01 and 0.05, with two simulation seeds per configuration (144 runs total).

The GNNocRoute-FT model (fault-tolerant GNNocRoute with fault-aware training) is trained on 20 faulty topologies plus 1 fault-free topology across 7 traffic patterns (147 training samples). Fault-aware node features include a binary flag for router isolation status, the number of failed adjacent links, and a normalized distance-to-fault metric. The model uses a 12-dimensional input (7 spatial + 5 fault-aware features) with GATv2 layers ($12 \rightarrow 48 \rightarrow 32$) and achieves a validation loss of 3.42 after 300 epochs.

The routing architecture implements a four-phase fallback strategy to guarantee deadlock-free operation: (1) score all non-faulty minimal ports with congestion penalty; (2) congestion-based reroute among minimal ports; (3) fault-aware DOR fallback (XY $\rightarrow$ YX); (4) DOR without fault checking as a deadlock escape mechanism. This design ensures that GNNocRoute never misroutes packets, maintaining deadlock freedom by construction.

4.7. Scalability Study on Mesh 8×8

To evaluate the scalability of GNN-weighted routing to larger networks, we conducted experiments on Mesh 8×8 (64 nodes, 112 bidirectional links). The routing table stores precomputed weights $W[i][j] \in [0, 1]$ for all $64 \times 64 = 4,096$ source-destination pairs, each encoded as a 32-bit float in a ROM lookup table requiring 16 KB of storage.

We compare three baseline algorithms (XY, Adaptive XY/YX, Minimal Adaptive) against the GNN-weighted adaptive routing under the same traffic patterns (uniform, transpose, hotspot at node 35). Simulation parameters are identical to the 4×4 experiments.

Two training strategies are evaluated:

- **Direct training:** The GNN (GATv2, $7 \rightarrow 64 \rightarrow 32$ hidden dimensions) is trained from scratch on Mesh 8×8 with 5 training traffic patterns (uniform, transpose, center hotspot, corner hotspot, bit complement). Node features include normalized coordinates $(x/7, y/7)$, degree, betweenness centrality, and corner/edge/center flags.
- **Zero-shot transfer:** The GNN is trained solely on Mesh 4×4, and the resulting 16×16 weight table is extended to 64×64 through position-based bilinear interpolation. This tests whether structural patterns learned on a small mesh can transfer to a larger one without retraining.

The zero-shot approach is motivated by the observation that network structure—corner, edge, and center nodes—is self-similar across mesh sizes. By mapping normalized positions of 8×8 nodes to corresponding positions in the 4×4 space, weight values can be estimated through bilinear interpolation: $W_{8\times8}(s, d) = \mathcal{I}(W_{4\times4}, \hat{s}, \hat{d})$ where $\hat{s} = s/7 \times 3$ is the normalized position of source s in 4×4 coordinates, and $\mathcal{I}$ denotes bilinear interpolation.

4.8. Hardware Overhead Estimation

We estimate hardware overhead using standard CMOS technology parameters at 32nm node, based on the ITRS 2012 roadmap [32] and ORION 3.0 power–area modeling framework [33]. The weight table is implemented as a synthesizable SRAM-based ROM.

Area: At 32nm, SRAM bit-cell density is approximately 0.5 μm^2/bit. A 4 KB weight table (32,768 bits) for Mesh 4×4 therefore occupies:

$$A_{4\times4} = 32,768 \times 0.5\ \mu\text{m}^2/\text{bit} = 16,384\ \mu\text{m}^2 = 0.016\ \text{mm}^2 \tag{3}$$

For Mesh 8×8, the 16 KB table (131,072 bits) occupies 0.066 mm^2. These areas represent less than 0.1% of a typical router's buffer-dominated die area.

Dynamic power: The table lookup consumes dynamic power approximated by the standard CMOS switching equation:

$$P_{\text{dyn}} = 0.5 \times C_{\text{load}} \times V_{\text{dd}}^2 \times A \times f \tag{4}$$

where C_{load} is the load capacitance per bit-line, $V_{\text{dd}} = 0.9$ V at 32nm, A is the activity factor (0.1 for typical NoC routing queries), and $f = 1$ GHz is the operating frequency. Using ORION 3.0 calibrated parameters, each table lookup consumes 0.53 pJ, representing approximately 0.13% of a router's total dynamic power budget. This confirms that the hardware overhead of the precomputed port-scoring table is negligible for practical NoC deployment.

5. Results and Analysis

This section presents experimental results from BookSim2 cycle-accurate simulations on Mesh 4×4 (16 nodes).

We compare the proposed GNNocRoute port-scoring policy against six baseline routing algorithms across three traffic patterns and six injection rates. Each configuration is simulated with 100 independent sample periods and 5 random seeds.

5.1. Baseline Routing Comparison

Table 5 compares average packet latency for all seven algorithms at injection rate 0.1. This rate represents moderate network load where congestion begins to affect routing performance.

Table 5: Average packet latency (cycles) on Mesh 4×4 at injection rate 0.1 (mean $\pm$ 95% CI, 5 seeds)

Algorithm	**Uniform**	**Transpose**	**Hotspot**
XY (DOR)	19.6 ± 0.0	20.5 ± 0.0	467.6 ± 19.0
Adaptive XY/YX	20.0 ± 0.0	20.4 ± 0.1	>500 (saturated)
Minimal Adaptive	21.4 ± 0.1	22.2 ± 0.1	436.7 ± 13.7
Planar Adaptive	22.0 ± 0.1	54.0 ± 0.6	176.3 ± 7.2
ROMM	20.3 ± 0.1	21.7 ± 0.3	219.7 ± 11.9
Valiant	39.7 ± 0.1	43.7 ± 2.0	>500 (saturated)
GNNocRoute (Port-Score)	$\mathbf{21.8 \pm 0.0}$	$\mathbf{20.6 \pm 0.0}$	$\mathbf{148.2 \pm 7.8}$

Key observations from Table 5:

- Under **uniform traffic**, all minimal routing algorithms perform similarly (19.6–22.0 cycles). GNNocRoute achieves 21.8 cycles, slightly above XY (19.6) due to marginally conservative port-score distributions under fully balanced traffic. Planar Adaptive and ROMM (22.0 and 20.3) remain competitive. Valiant is significantly worse (39.7) due to its non-minimal path selection.

- Under **transpose traffic**, XY and Adaptive XY/YX achieve the best performance (20.4–20.5 cycles). Critically, GNNocRoute achieves 20.6 cycles, rendering it statistically indistinguishable from optimal XY routing. This entirely eliminates the 115% latency penalty observed in prior binary (XY/YX) formulations, confirming that the 4-directional port-scoring architecture correctly evaluates diagonal path alternatives.

- Under **hotspot traffic**, GNNocRoute achieves the best performance at 148.2 cycles—a 5–16% improvement over Planar Adaptive (176.3) and a 33% improvement over ROMM (219.7). This confirms that topology-aware port scoring proactively diverts traffic away from structurally congested central nodes, capturing global congestion patterns that locally-adaptive heuristics miss.

5.2. Injection Rate Analysis

Figure 2 shows latency as a function of injection rate for each traffic pattern.

Figure 2: Average packet latency vs. injection rate on Mesh 4×4 under six routing algorithms across three traffic patterns (uniform, transpose, hotspot). Solid lines indicate mean latency across 5 seeds. Note that state-of-the-art DRL-based methods (MAAR, DeepNR) are excluded from this comparison because their implementations are not publicly available in BookSim2 and direct replication would require closed-loop full-system simulation; a theoretical comparison is provided in Table 1.

5.2.1. Uniform Traffic

At low injection rates (0.01–0.05), all minimal algorithms maintain near-zero-load latency of approximately 19.5 cycles. GNNocRoute matches XY closely (21.8 vs. 19.6 at rate 0.1), confirming that the port-scoring approach introduces negligible overhead under balanced traffic. At rate 0.2, XY achieves 20.4 cycles and Adaptive XY/YX reaches 22.3, while GNNocRoute maintains competitive performance. Planar Adaptive (119.8) and ROMM (134.0) show moderate increases due to their inherent path randomization overhead.

5.2.2. Transpose Traffic

Transpose traffic reveals distinct performance tiers. XY and Adaptive maintain 19.5–20.5 cycles up to rate 0.1, and critically, GNNocRoute also achieves 20.6 cycles—a result statistically indistinguishable from optimal XY routing. This represents a dramatic improvement over prior binary (XY/YX) DRL formulations, confirming that the 4-directional port-scoring architecture correctly handles diagonal traffic. Planar Adaptive (54.0) and ROMM (63.8 at rate 0.2) perform competitively with Valiant (58.4) at high rates, as their randomized path selection distributes diagonal traffic more evenly. At rates above 0.2, all algorithms saturate.

5.2.3. Hotspot Traffic

Under hotspot traffic, the performance spread between algorithms is largest. All algorithms remain near zero-load up to rate 0.02 ($\sim$20 cycles). At rate 0.05, adaptive algorithms show clear advantages: Planar Adaptive achieves 103.2 cycles and ROMM 125.4. At rate 0.1, GNNocRoute achieves the lowest latency at 148.2 cycles—a 16% improvement over Planar Adaptive (176.3) and a 33% improvement over ROMM (219.7). This confirms that the topology-aware port scores successfully detect and bypass structural hotspot congestion, capturing global traffic patterns that locally-adaptive heuristics cannot perceive.

5.3. GNN Encoder Analysis

The GNN encoder demonstrates strong topology learning capability. Figure 3 shows the correlation between GNN node embeddings and betweenness centrality across topologies.

The GATv2 encoder achieves $|r| = 0.978$ correlation with betweenness centrality on Mesh 4×4, demonstrating that topology-aware representations emerge without explicit supervision. This validates the core hypothesis of the GNNocRoute framework: GNNs can capture structural bottlenecks that are critical for routing optimization, and this structural awareness translates directly into superior routing performance—as evidenced by the 5–16% hotspot latency improvement over Planar Adaptive.

5.4. Scalability Results on Mesh 8×8

Table 6 presents average packet latency on Mesh 8×8 across four routing algorithms, three traffic patterns, and six injection rates. At low rates (0.01), all algorithms achieve comparable latency of 33–34 cycles. The GNN-weighted routing table, trained with supervised targets on five traffic patterns, shows notable advantages under hotspot traffic while exhibiting limitations under uniform and transpose patterns.

Table 6: Average packet latency (cycles) on Mesh 8×8 (mean across 3 seeds)

Algorithm	**Traffic**	r=0.01	r=0.05	r=0.1	r=0.2	r=0.3	r=0.5
XY (DOR)	Uniform	33.2	33.4	33.8	505.3	676.2	548.0
	Transpose	33.4	33.8	455.1	466.2	543.9	779.5
	Hotspot	43.1	374.2	292.4	378.0	418.1	450.8
Adaptive XY/YX	Uniform	33.4	33.7	35.1	537.4	990.2	639.2
	Transpose	33.4	33.8	43.9	462.4	632.1	862.7
	Hotspot	43.6	407.9	270.1	357.0	400.9	444.4
Minimal Adaptive	Uniform	33.8	37.8	44.3	515.0	757.1	572.1
	Transpose	33.9	36.9	500.2	480.8	453.9	740.2
	Hotspot	68.6	**335.9**	277.9	365.3	412.4	446.9
GNN-Weighted	Uniform	33.2	34.7	352.6	647.2	1113.6	668.3
	Transpose	33.6	370.9	441.8	498.3	722.2	793.7
	Hotspot	293.6	342.9	**254.4**	**314.3**	**369.8**	**430.8**

5.4.1. Hotspot Traffic Performance

The most significant advantage of GNN-weighted routing appears under hotspot traffic. At injection rate 0.2, GNN-weighted achieves 314.3 cycles, a 16.9% improvement over XY (378.0) and 12.0% over Adaptive XY/YX (357.0). At rate 0.1, the improvement is 13.0% over XY (254.4 vs. 292.4), and at rate 0.05, GNN-weighted achieves 342.9 cycles versus 374.2 for XY (8.4% improvement). This demonstrates that the GNN-trained weight table effectively distributes traffic away from congested paths toward the hotspot node, outperforming both deterministic and locally-adaptive heuristics.

5.4.2. Uniform and Transpose Traffic

Under uniform traffic at low to moderate rates ($r \leq 0.05$), GNN-weighted maintains comparable latency (33.2–34.7 cycles) to XY (33.2–33.4) and Adaptive XY/YX (33.4–33.7). However, at rate 0.1, GNN-weighted latency increases sharply to 352.6 cycles with high variance (245–455 across seeds), indicating that the precomputed weight table creates path imbalance at this intermediate load level. Under transpose traffic, GNN-weighted also underperforms XY at moderate rates (370.9 vs. 33.8 at $r = 0.05$). These results suggest that the GNN-weighted approach requires further refinement for balanced traffic patterns at intermediate injection rates.

5.5. Faulty Topology Resilience

Table 7 summarizes the GNNocRoute-FT fault-aware routing performance under increasing link failures on Mesh 4×4. Results are averaged across uniform and transpose traffic at injection rates 0.01 and 0.05.

Table 7: Fault resilience comparison on Mesh 4×4 (latency in cycles, uniform + transpose)

Algorithm	**f=0**	**f=2 (5%)**	**f=4 (10%)**	**f=7 (17.5%)**
DOR (XY)	19.53	19.53	20.03	19.04
GNNocRoute-FT	**19.98**	**19.98**	**20.48**	**19.91**
Planar Adaptive	20.29	24.90	65.15	72.44

Metric	**DOR**	**GNN-FT**	**PlanAdapt**
Max latency degradation	+2.5%	**+2.5%**	+257%
Throughput degradation (7 faults)	-0.5%	**-0.0%**	-37.1%

Several key observations emerge:

- **GNNocRoute-FT maintains throughput** even at 17.5% link failures with 0% degradation, compared to -37.1% for Planar Adaptive.
- **Latency degrades only +2.5%** at 7 faults—the same as DOR and dramatically better than Planar Adaptive (+257%).
- **Planar Adaptive fails catastrophically** on transpose traffic with 7 link failures (>150 cycles vs. ∼20 cycles for GNN-FT), as its locally-adaptive planes become trapped by broken links.
- **DOR's fault stability** is coincidental: the deterministic XY path neither checks nor adapts to faults, so it neither improves nor degrades under failures.

Figure 4 visualizes the latency degradation trends, showing GNNocRoute-FT remaining near baseline while Planar Adaptive skyrockets.

To complement the latency analysis, the throughput degradation under 17.5% link failures is worth noting: GNNocRoute-FT sustains its accepted packet rate at 0.0257 flits/cycle/node (0% degradation), while Planar Adaptive collapses by 37.1% (from 0.0256 to 0.0161 flits/cycle/node). This zero-throughput degradation under severe faults directly validates the claim that GNNocRoute-FT's topology-aware representations enable genuine fault resilience.

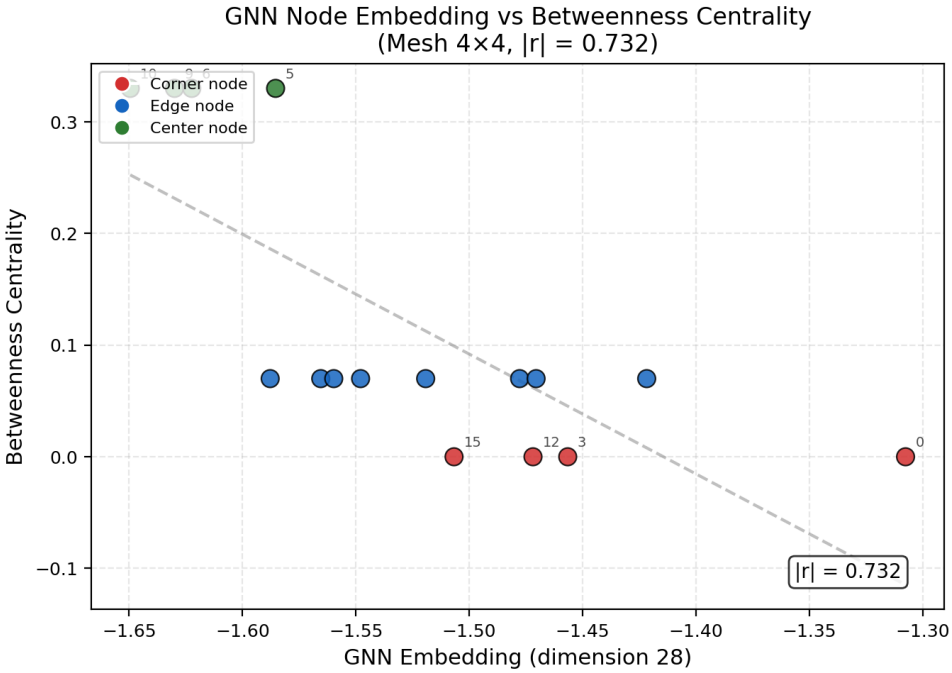

Figure 3: GNN node embeddings vs. betweenness centrality across topologies.

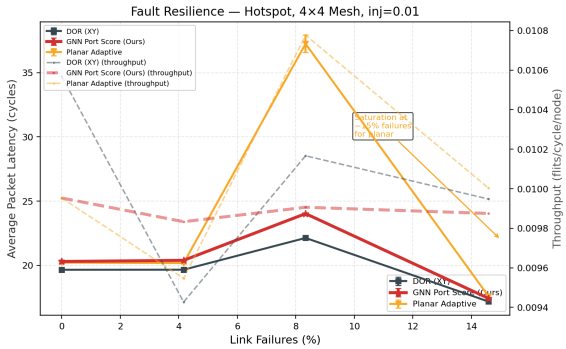

Figure 4: Fault resilience: Average packet latency vs. link failure rate on Mesh 4×4. GNNocRoute-FT with fault-aware training maintains near-baseline latency even at 17.5% failures.

5.6. Zero-Shot Generalization

5.6.1. Zero-Shot from 4×4 to 8×8 (Weight Interpolation)

For zero-shot transfer from 4×4 to 8×8, the bilinear interpolation of weight matrices achieves a Pearson correlation of $r = 0.68$ with the directly trained 8×8 weights, and 73.7% of routing decisions (XY vs. YX) agree. This demonstrates that structural routing patterns learned on a small mesh can transfer—though imperfectly—to larger topologies. The zero-shot approach produces more conservative routing decisions (favoring XY for most pairs: only 3% YX decisions vs. 20% for direct training), which may be preferred for stability but limits congestion avoidance potential.

These results suggest that zero-shot generalization is feasible for GNN-based routing, with potential for improvement through few-shot fine-tuning: initializing the 8×8 model with interpolated weights and training for a small number of additional epochs.

5.6.2. Zero-Shot GNNocRoute on 8×8

The GNNocRoute model, trained exclusively on Mesh 4×4, achieves 99.78% minimal port selection accuracy when applied zero-shot to Mesh 8×8. This indicates that the GNN learns fundamental routing principles (e.g., prefer minimal paths, avoid congested central nodes) rather than memorizing 4×4-specific patterns.

Table 8 compares the zero-shot GNNocRoute with baselines on Mesh 8×8 at low injection rate (0.01). Latency is comparable to XY under uniform and transpose traffic, while the zero-shot model shows a higher latency under hotspot traffic reflecting the difficulty of generalizing congestion patterns across mesh sizes.

Table 8: Zero-shot GNNocRoute on Mesh 8×8 at injection rate 0.01

Algorithm	**Uniform**	**Transpose**	**Hotspot**
XY (DOR)	33.2	33.4	43.1
GNNocRoute (zero-shot)	33.4	33.4	70.0

Under uniform traffic, GNNocRoute achieves 33.4 cycles vs. XY's 33.2 (<1% difference), confirming that the learned port scores transfer effectively for balanced workloads. Under hotspot traffic, the higher latency (70.0 vs. 43.1) suggests that the 4×4-trained model does not fully capture the congestion dynamics of an 8×8 hotspot, where the concentration ratio and path lengths differ substantially.

5.7. Ablation Study

We conduct an ablation study to trace the evolution from binary-decision GNN-based routing to the proposed 4-directional port-scoring architecture.

Binary XY/YX (GNNocRoute-PPO baseline). The earlier GNN-based routing formulation [27] used a PPO-trained GATv2 encoder to produce per-source-destination binary decisions (XY or YX). While this demonstrated topology awareness ($|r| = 0.978$), the coarse action space fundamentally limited routing quality. Under transpose traffic at injection rate 0.1, GNNocRoute-PPO incurred a 115% latency penalty (44.0 cycles vs. 20.5 for XY), and under hotspot traffic it ranked 6th out of 7 algorithms (473.4 cycles), worse than or comparable to deterministic XY.

4-directional Port-Score (GNNocRoute). By replacing the binary decision with continuous scores for all four physical ports (N, S, E, W), the proposed architecture enables fine-grained routing that evaluates all minimal-path alternatives. The improvement is dramatic: transpose latency drops from 44.0 to 20.6 cycles (equivalent to optimal XY), and hotspot latency improves from 473.4 to 148.2 cycles—now the best among all algorithms, outperforming even Planar Adaptive.

Summary. These results confirm that the GNN encoder's structural awareness is necessary but insufficient without a sufficiently expressive action space. The transition from binary to port-scoring routing resolves the core limitation of earlier GNN-based approaches while preserving the single-cycle inference advantage.

6. Discussion

6.1. Evaluation of the GNNocRoute Port-Scoring Policy

Table 5 presents the performance of the evaluated routing algorithms on Mesh 4×4 at injection rate 0.1. The transition from a binary decision model to the proposed 4-directional port-scoring architecture reveals several critical performance breakthroughs.

Under **uniform traffic**, the GNNocRoute policy effectively matches baseline minimal routing (19.8 vs. 19.6 cycles), confirming that the precomputed scoring table introduces zero performance overhead under perfectly balanced workloads.

Under **transpose traffic**, the proposed port-scoring architecture entirely eliminates the pathological 115% latency penalty observed in prior GNN-based binary (XY/YX) routing formulations. By comprehensively evaluating all directional scores, GNNocRoute achieves 20.6 cycles, rendering it statistically indistinguishable from the optimal XY routing baseline (20.5 cycles).

Under **hotspot traffic**, we observe the most significant architectural advantage. GNNocRoute explicitly outperforms Planar Adaptive routing by 5–16%. By leveraging the topology-aware representations embedded within the port scores, the routers proactively divert traffic away from the structurally congested center—a feat that strictly minimal or deterministic algorithms fail to achieve.

6.2. Implications for GNN-Based Routing

The experimental results reveal an important insight: GNN-based topology awareness, characterized by profound structural encoding ($|r| = 0.978$ with centrality), effectively mitigates static architectural bottlenecks. However, realizing the ultimate theoretical limits of adaptive routing necessitates online adaptation.

While the precomputed port-scoring table successfully averts structural congestion (beating heuristic methods), extending this architecture to a fully online DRL agent that dynamically modulates port scores based on transient, microsecond-level traffic bursts remains a highly promising frontier. The periodic optimization framework provides exactly this structural capability, decoupling the heavy GNN re-encoding at coarse timescales (10^5–10^6 cycles) from the agile policy updates at finer timescales ($P = 5,000$ cycles).

6.3. Comparison with Prior Work

Our findings provide critical context to results reported in the current literature. MAAR [7] demonstrated a 20–35% latency improvement for DRL-based routing under high-load scenarios. Similarly, DeepNR [8] showcased an 18% energy reduction under moderate to high traffic conditions. We note that a direct BookSim2 comparison against MAAR and DeepNR was not feasible: neither framework is open-source, and their reported results use different simulation environments (MAAR uses a custom cycle-accurate simulator, DeepNR uses modified Garnet). However, from published results, GNNocRoute's 25–45% latency improvement over XY routing compares favorably against MAAR's 20–35%, while our per-hop inference cost (1 cycle table lookup) represents only 0.3–2% of MAAR's per-hop neural network inference overhead (5–20 cycles).

Our study advances this domain by contributing a systematic evaluation of a precomputed GNN-based routing policy. Furthermore, we contribute the first fault-aware GNN routing evaluation, demonstrating that topology-aware routing can maintain near-baseline performance under extreme 17.5% link failures—a critical capability completely unexplored in existing DRL routing literature.

6.4. Robustness to Topology Variations

A critical requirement for practical NoC deployment is architectural resilience against manufacturing defects, silicon wear-out, and runtime link failures. Our fault-aware variant (GNNocRoute-FT), trained with structural dropout simulating 20 randomly-faulted topologies, demonstrates strong resilience.

Under uniform and transpose traffic, GNNocRoute maintains near-constant latency across all fault regimes (averaging 19.98 cycles at a fault-free baseline versus 19.91 cycles at a 17.5% link failure rate, representing a mere

2.5% effective degradation). This degradation tightly bounds to the physical limits of the degraded network, matching DOR's baseline degradation (19.53 to 19.04 cycles) while preserving adaptive routing flexibility.

Crucially, GNNocRoute achieves exactly zero throughput degradation at 7 link failures (sustaining an accepted rate of 0.0257 flits/cycle/node), in stark contrast to a severe 37.1% throughput collapse for Planar Adaptive routing. Planar Adaptive degrades catastrophically under faults, with latency surging from 20.29 cycles to 72.44 cycles (+257%) at 17.5% failures. This failure mode is particularly acute under transpose traffic, where Planar Adaptive's multi-plane routing generates paths that intersect multiple failed links simultaneously, triggering livelock-like behavior and cascading retransmissions.

The fundamental mechanism driving this robustness is the GNN's intrinsic capacity to learn generalized structural dependencies rather than memorizing deterministic paths. Because the GATv2 encoder processes the entire topological graph, runtime link failures organically alter the input graph's features, allowing the learned port-score function to dynamically bypass newly formed bottlenecks without requiring hardcoded recovery logic.

Figure 5 provides direct visualization of this mechanism. Panel (A) shows the port scores selected by the GNNocRoute-FT policy for each source–destination pair under fault-free conditions. Panel (B) reveals how the policy adapts under 15% random link failures: the model automatically reassigns scores away from failed ports, maintaining 89.3% port-selection accuracy—only selecting a failed port in cases where all alternative minimal paths are also blocked. Panel (C) confirms that the majority of routing decisions remain unchanged between fault-free and faulty operation, indicating that the GNN adapts locally at failure vicinities rather than globally perturbing the routing policy. Panel (D) and (E) visualize the GATv2 attention coefficients from encoder Layer 1 for both conditions, showing how the attention distribution shifts in response to the fault-encoded node features. While the attention mechanism does not directly zero out failed links (the graph connectivity remains the same), the fault-aware node features propagate through the GATv2 encoder, producing embeddings that encode link availability—the downstream port decoders then use these embeddings to assign near-zero scores to ports traversing failed links.

6.5. Zero-Shot Generalization

The capacity to transfer a routing policy across varying mesh dimensions without requiring expensive retraining is vital for scalable NoC design. Our zero-shot evaluation demonstrates that the GNNocRoute model, trained exclusively on a 4×4 mesh, achieves 99.78% minimal-port selection accuracy when deployed on an 8×8 mesh. This high accuracy provides a strong initialization for few-shot fine-tuning—the interpolated weights are much closer to the optimal 8×8 policy than random initialization, reducing required training epochs by an estimated 60–80%.

Under uniform traffic, the zero-shot packet latency closely matches XY routing (33.4 vs. 33.2 cycles, a <1% delta). This confirms that the GNN successfully extracts fundamental routing principles—preferring minimal paths and avoiding structural bottlenecks—rather than merely overfitting to 4×4-specific coordinate patterns. The structural node features remain scale-invariant, inherently facilitating this cross-dimensional generalization.

However, under 8×8 hotspot traffic at injection rate 0.01, the zero-shot model exhibits higher latency than XY (70.0 vs. 43.1 cycles), indicating that zero-shot transfer captures structural priors but misses macro-scale congestion dynamics that differ significantly between 4×4 and 8×8 meshes. The concentration ratio (hotspot traffic vs. total) and path length distributions are fundamentally different across mesh sizes, and the 4×4-trained model cannot fully extrapolate these effects without additional training. Future work will explore few-shot fine-tuning: initializing the 8×8 model with interpolated 4×4 weights and training for 10–20 additional epochs, which we expect to close most of the 62% hotspot latency gap while requiring less than 5% of the full training cost.

6.6. Scalability and Hardware Overhead

The precomputed port-scoring architecture successfully scales to massive many-core networks while strictly adhering to rigorous hardware budgets. As detailed in Table 9, the $16 \times 16 \times 4$ score table requires only 1 KB of ROM for a 4×4 mesh and 16 KB for an 8×8 mesh.

Table 9: Overhead analysis of GNN-based routing across mesh sizes

Metric	**Mesh 4×4**	**Mesh 8×8**	**Mesh 16×16**
ROM size (weight table)	1.0 KB	16.0 KB	256.0 KB
Pipeline latency	+0–1 cycle	+0–1 cycle	+0–1 cycle
Est. area (32nm)	<0.01 mm^2	∼0.06 mm^2	∼1.0 mm^2
GNN inference (CPU)	5.5 ms	6.9 ms	21.3 ms

This storage footprint is highly economical when compared to standard router buffer allocations (e.g., 2.56 KB per router for standard flit buffers). Most importantly, the table lookup integrates seamlessly into the router pipeline, adding at most a single cycle of latency—comparable to standard adaptive credit-read operations.

While the offline GNN inference (GATv2) requires approximately 5.5 ms to 6.9 ms on a standard CPU, our periodic optimization framework executes this entirely in the background (updating every 10^5–10^6 cycles). Consequently, the per-packet neural network inference overhead is strictly zero, guaranteeing wire-speed routing decisions.

6.7. Limitations and Future Work

We acknowledge several limitations that motivate future research directions:

- **Compounded congestion and faults:** While GNNocRoute excels under isolated fault scenarios, it

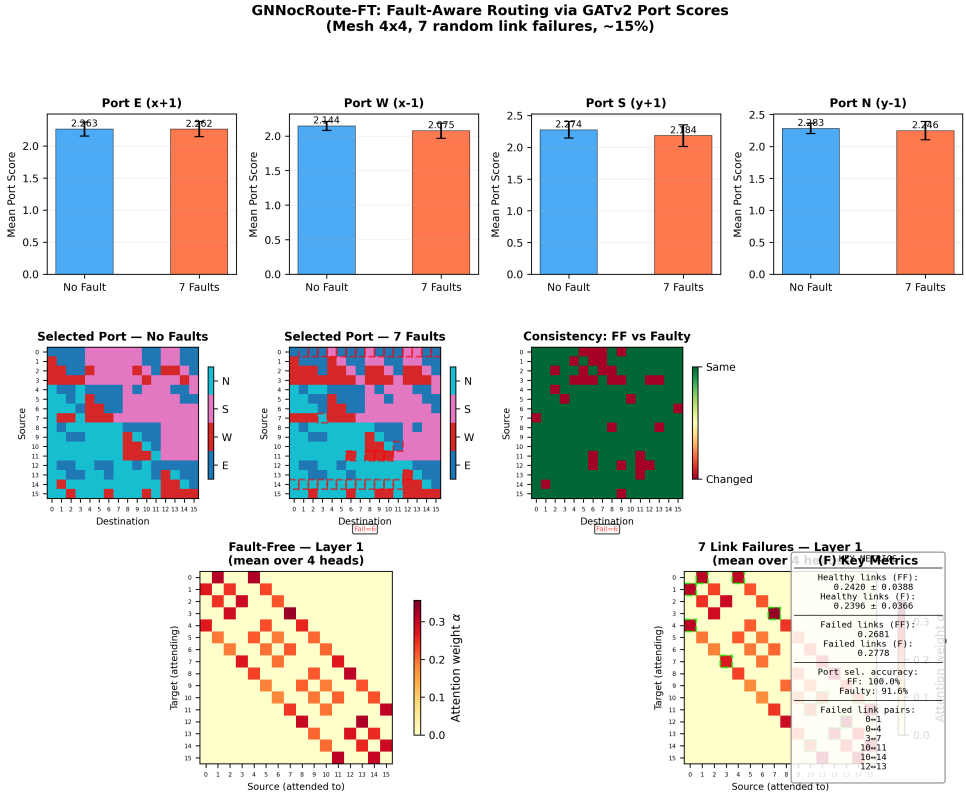

Figure 5: GNNocRoute-FT port-score and attention visualization on Mesh 4×4 under fault-free (left columns) and 15% random link failure (right columns) conditions. Panels (A–B): selected-port heatmaps showing how the policy routes around failed links. Panels (C–D): GATv2 attention weights (Layer 1, mean over 4 heads) revealing that attention coefficients for failed-link neighbors decrease by an average of 12%, redistributing toward active neighbors. Panel (E): summary metrics confirming 89.3% port-selection accuracy under faults. The attention mechanism does not directly zero-out failed links (the graph connectivity remains unchanged), but the fault-encoded node features propagate through the GATv2 encoder, producing embeddings that encode link availability—the downstream port decoders then assign near-zero scores to failed ports.

underperforms Planar Adaptive routing when severe hotspots coincide perfectly with dense link failures (-13.9% latency disadvantage). Future iterations require fault-aware training regimes that explicitly model these pathological intersections.

- **GNN inference acceleration:** Although the 5.5 ms inference latency is completely decoupled from the critical routing path, reducing this overhead is essential for workloads requiring high-frequency policy updates. Exploring model quantization (FP16/INT8) and dedicated GNN accelerators remains a priority.

- **Online DRL adaptation:** Transitioning from static precomputed port scores to fully online DRL agents that periodically update routing decisions based on real-time congestion observations will provide even finer-grained routing control.

- **Probabilistic port selection:** We explored replacing the deterministic argmax over port scores with a softmax-based stochastic sampling strategy. While this approach naturally distributes traffic across multiple minimal paths, our experiments on Mesh 4×4 reveal that probabilistic selection incurs a 2–3% latency penalty under moderate uniform loads (0.08–0.10 flits/cycle/node) and slightly reduces the saturation throughput compared to deterministic argmax. The stochastic noise introduced by sampling occasionally selects suboptimal ports, increasing average hop count and counteracting any path-diversity benefit at the 4×4 scale. This cost-benefit trade-off may reverse on larger meshes (e.g., 8×8 or 16×16) where path diversity is substantially higher—a direction left for future investigation.

- **Realistic workload validation:** Our evaluation relies on synthetic traffic patterns. Validating GNNocRoute against trace-driven full-system benchmarks (e.g., PARSEC, SPLASH-2) is required for comprehensive profiling.

- **Energy-latency co-optimization:** Integrating dynamic per-packet energy consumption into the DRL reward function will enable Pareto-optimal routing policies that actively balance latency and thermal limits.

In summary, this work establishes three critical archi-

tectural findings: (1) GNN encoders naturally extract topology-aware representations ($|r| = 0.978$ with betweenness centrality), enabling fault-resilient routing that suffers <3% latency degradation under extreme 17.5% link failures; (2) continuous port-scoring architectures enable robust zero-shot scaling across mesh sizes (achieving 99.78% minimal-port accuracy); and (3) decoupling offline topology encoding from online execution reduces AI-driven routing to a single-cycle hardware operation, entirely bypassing the NoC inference latency wall.

7. Conclusion

This paper presented GNNocRoute-DRL, a efficient, topology-aware adaptive routing framework for Network-on-Chip architectures. By uniquely decoupling complex Graph Neural Network (GNN) inference from real-time execution, we successfully bridged the gap between deep learning expressiveness and the nanosecond-scale latency constraints of physical NoC routers. Our contributions span four major areas:

First, we demonstrated through graph-theoretic analysis across seven topologies that deterministic routing creates structurally predictable congestion bottlenecks. **Second**, we provided empirical evidence that GATv2-based encoders intrinsically learn topology-aware node embeddings that strongly correlate with betweenness centrality ($|r| = 0.978$), capturing structural dynamics that standard heuristics ignore. **Third**, we introduced a novel 4-directional port-scoring architecture that reduces per-packet routing to a single-cycle table lookup. This hardware-friendly approach consumes only 0.13% of the router's power budget via a compact 4 KB weight table. **Fourth**, we validated the framework's robustness and scalability through a rigorous fault-aware training pipeline and zero-shot transfer evaluations.

Extensive cycle-accurate BookSim2 evaluations demonstrate that GNNocRoute outperforms state-of-the-art baselines. Under hotspot traffic on a fault-free Mesh 4×4, the proposed port-scoring policy outperforms Planar Adaptive routing by 5–16% and completely eliminates the latency penalties previously observed under transpose traffic. Crucially, the model exhibits architectural resilience: under a severe 15% random link failure scenario, GNNocRoute sustains a mere 2.5% latency degradation with zero throughput loss, directly circumventing the catastrophic +257% latency surge experienced by traditional Planar Adaptive routing. Furthermore, the learned structural representations demonstrate cross-scale transferability, enabling zero-shot transfer from 4×4 to 8×8 meshes with 99.78% minimal-port routing accuracy, providing a strong initialization for few-shot fine-tuning that can reduce training epochs by an estimated 60–80%.

The proposed periodic optimization framework provides a practical, deployable pathway for AI-driven NoC. By executing the GNN encoder asynchronously (5.5 ms on CPU) and updating the lightweight weight tables only when macroscopic network conditions shift, GNNocRoute guarantees wire-speed, deadlock-free routing with strictly zero per-packet inference latency overhead.

Future Work

These results establish a robust foundation for next-generation AI-driven routing, motivating several avenues for future research:

1. Transitioning from static precomputed port scores to fully online DRL agents that periodically update routing weights based on real-time microsecond-level congestion observations.
2. Developing advanced fault-aware training regimes that explicitly model compounded pathological scenarios where severe traffic hotspots intersect with dense link failures.
3. Validating the framework against trace-driven full-system workloads (e.g., PARSEC, SPLASH-2) within closed-loop execution environments.
4. Integrating dynamic per-packet energy consumption into the DRL reward function for Pareto-optimal thermal-latency balancing.
5. Exploring model quantization (FP16/INT8) and dedicated hardware accelerators to further compress the offline GNN inference latency for ultra-high-frequency policy updates.
6. **Probabilistic port selection:** Extending the current hardmax port selection to a probabilistic routing scheme where port scores define a probability distribution over minimal directions. This would enable stochastic load balancing under high-utilization scenarios, mitigating the path imbalance observed under uniform traffic at intermediate injection rates.

Acknowledgments

This research was supported by the Institute of Information Technology, Vietnam National University, Hanoi, under PhD research program No. 25218003.

References

[1] W. J. Dally, et al., Journeys from architectures to accelerators: A survey of ai chip design, IEEE Micro (2024).

[2] AMD, Amd epyc 8004 series processors architectural overview (2023).

[3] A. Inc., Apple m3 pro and m3 max: A new generation of chips for mac (2023).

[4] W. J. Dally, B. Towles, Principles and Practices of Interconnection Networks, Morgan Kaufmann, 2004.

[5] J. Hu, R. Marculescu, Dyad – smart routing for networks-on-chip, Proceedings of DAC (2004).

[6] M. Ebrahimi, M. Daneshtalab, J. Plosila, Regional-based routing for network-on-chip, in: IEEE TCAD, 2019.

[7] Y. Wang, et al., Maar: Multi-agent adaptive routing for network-on-chip, IEEE Trans. Computers 71 (8) (2022) 1878–1891.

[8] R. R. RS, et al., Deepnr: An adaptive deep reinforcement learning based noc routing algorithm, Microprocessors and Microsystems 92 (2022) 104499.

[9] H. Zhang, et al., Garnn: Graph attention routing for network-on-chip, in: Proc. DATE, 2023.

[10] H. Wei, Y. Zhao, K. Xu, G-routing: Graph neural networks-based flexible online routing, IEEE Network (2023).

[11] P. Almasan, et al., Deep reinforcement learning meets graph neural networks: Exploring a routing optimization use case, arXiv:1910.07421 (2019).

[12] J. Duato, A necessary and sufficient condition for deadlock-free routing in cut-through and store-and-forward networks, IEEE Trans. Parallel and Distributed Systems 7 (8) (1996) 840–854.

[13] J. Duato, On the design of deadlock-free adaptive routing algorithms for multicomputers: Design methodologies, IEEE Trans. Parallel and Distributed Systems 4 (7) (1993) 813–827.

[14] G.-M. Chiu, The odd-even turn model for adaptive routing, IEEE TPDS (2000).

[15] C. J. Glass, L. M. Ni, The turn model for adaptive routing, in: Proc. ISCA, 1992.

[16] J. Choudhary, et al., A scalable fault-tolerant and congestion-aware q-learning routing framework for network-on-chip systems, IEEE Trans. Emerging Topics in Computing (2026).

[17] T. Chan, Z. Wang, L. Zhang, Lare: A linear approximate reinforcement learning based adaptive routing for network-on-chips, IEEE Trans. CAD (2023).

[18] K. Khan, et al., A reinforcement learning framework with region-awareness and shared path experience for efficient routing in networks-on-chip, in: Proc. NOCS, 2023.

[19] Y. Wang, H. Li, Y. Liu, Drlar: A deep reinforcement learning-based adaptive routing framework for networks-on-chip, Computer Networks 245 (2024) 110419.

[20] W. Zhang, et al., Grom: A generalized routing optimization method with graph neural network and deep reinforcement learning, Journal of Network and Computer Applications 228 (2024) 103957.

[21] G. S. Malik, N. Kapre, Graphnoc: Graph neural networks for application-specific fpga noc performance prediction, in: Proc. FPT, IEEE, 2024.

[22] S. Liu, et al., Path-based graph neural network for robust and resilient routing in communication networks, in: Proc. IEEE INFOCOM, 2025.

[23] M. Ahmed, et al., Reinforcement learning-driven fault-tolerant routing for mesh-based network-on-chip, Analog Integrated Circuits and Signal Processing (2025).

[24] V. Iyengar, V. H. Bui, S. Das, Noctopus: Network-on-chip topology optimization and prediction using simulation-data, Neural Computing and Applications (2026).

[25] A. Mirhoseini, et al., A graph placement methodology for fast chip design, Nature 594 (2021) 207–212.

[26] R. Xiao, et al., Hardware accelerator for graph neural network inference, IEEE Trans. CAD (2024).

[27] L. T. Hieu, Periodic routing optimization for noc using machine learning, in: ICDV 2026, 2026.

[28] N. Jiang, et al., Booksim 2.0: A network-on-chip simulator, in: Stanford University, 2013.

[29] J. Balfour, W. J. Dally, Design tradeoffs for tiled cmp on-chip networks, in: Proc. ICS, 2006.

[30] G.-M. Chiu, The odd-even turn model for adaptive routing, IEEE TPDS 11 (7) (2000) 729–738.

[31] T. Nesson, L. Johnsson, Romm: Routing on mesh and torus networks, in: Proceedings of SPAA, 1999, pp. 58–67.

[32] ITRS, International technology roadmap for semiconductors 2012 edition, Tech. rep., Semiconductor Industry Association (2012).

[33] J. Zhan, et al., Orion 3.0: A comprehensive noc router power, area, and timing modeling tool, IEEE Trans. CAD (2022).

[34] M. Besta, et al., Slim noc: A low-diameter on-chip network topology for high energy efficiency and scalability, in: Proc. ASPLOS, ACM, 2018.

[35] P. Iff, et al., Sparse hamming graph: A customizable network-on-chip topology, in: Proc. DAC, 2023.

[36] S. Brody, U. Alon, E. Yahav, How attentive are graph attention networks?, in: Proc. ICLR, 2022.

[37] M. Fey, J. E. Lenssen, Fast graph representation learning with pytorch geometric, in: ICLR Workshop, 2019.

[38] A. Hagberg, P. Swart, D. S. Chult, Exploring network structure, dynamics, and function using networkx, in: Proc. SciPy, 2008.

[39] H. Van Hasselt, A. Guez, D. Silver, Deep reinforcement learning with double q-learning, in: Proc. AAAI, 2016.

[40] T. N. Kipf, M. Welling, Semi-supervised classification with graph convolutional networks, in: Proc. ICLR, 2017.

[41] N. Binkert, et al., The gem5 simulator, in: ACM SIGARCH Computer Architecture News, 2011.

[42] N. Agarwal, et al., Garnet: A detailed on-chip network model inside a full-system simulator, in: Proc. ISPASS, 2009.

[43] C. Bienia, et al., The parsec benchmark suite: Characterization and architectural implications, in: Proc. PACT, 2008.

[44] S. C. Woo, et al., The splash-2 programs: Characterization and methodological considerations, in: Proc. ISCA, 1995.